

Class 08: Breast Cancer Mini Project

Erin Jang (PID: A17713992)

Table of contents

Background	1
Data Import	1
Exploratory data analysis	3
Principal Component Analysis (PCA)	4
Variance explained	7
Communicating PCA results	10
Hierarchical clustering	10
Combining methods	12
Prediction	15
About this document	16

Background

The goal of this mini-project is for you to explore a complete analysis using the unsupervised learning techniques covered in class.

Today we will analyze a data-set from fine needle aspiration (FNA) of a breast mass.

Data Import

The data is made available as a CSV file for download. We can read this in using `read.csv()`:

```
wisc.df <- read.csv("WisconsinCancer.csv", row.names = 1)
```

```
head(wisc.df, 3)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	
842302	M	17.99	10.38	122.8	1001	
842517	M	20.57	17.77	132.9	1326	
84300903	M	19.69	21.25	130.0	1203	
		smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302		0.11840	0.27760	0.3001	0.14710	
842517		0.08474	0.07864	0.0869	0.07017	
84300903		0.10960	0.15990	0.1974	0.12790	
		symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302		0.2419	0.07871	1.0950	0.9053	8.589
842517		0.1812	0.05667	0.5435	0.7339	3.398
84300903		0.2069	0.05999	0.7456	0.7869	4.585
		area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302		153.40	0.006399	0.04904	0.05373	0.01587
842517		74.08	0.005225	0.01308	0.01860	0.01340
84300903		94.03	0.006150	0.04006	0.03832	0.02058
		symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
842302		0.03003	0.006193	25.38	17.33	
842517		0.01389	0.003532	24.99	23.41	
84300903		0.02250	0.004571	23.57	25.53	
		perimeter_worst	area_worst	smoothness_worst	compactness_worst	
842302		184.6	2019	0.1622	0.6656	
842517		158.8	1956	0.1238	0.1866	
84300903		152.5	1709	0.1444	0.4245	
		concavity_worst	concave.points_worst	symmetry_worst		
842302		0.7119	0.2654	0.4601		
842517		0.2416	0.1860	0.2750		
84300903		0.4504	0.2430	0.3613		
		fractal_dimension_worst				
842302		0.11890				
842517		0.08902				
84300903		0.08758				

Make sure we remove or exclude the **diagnosis** column from the data-set that we use for further analysis - this is the expert diagnosis as either M or B.

```
# We can use -1 here to remove the first column
wisc.data <- wisc.df[,-1]
diagnosis <- as.factor( wisc.df$diagnosis )
```

Exploratory data analysis

Q1. How many observations are in this dataset?

```
nrow(wisc.data)
```

```
[1] 569
```

There are 569 observations.

Q2. How many of the observations have a malignant diagnosis?

```
table(wisc.df$diagnosis)
```

```
  B   M  
357 212
```

212 observations have a malignant diagnosis.

Q3. How many variables/features in the data are suffixed with `_mean`?

We can use the `grep()` function to help us here:

```
colnames(wisc.data)
```

```
[1] "radius_mean"      "texture_mean"  
[3] "perimeter_mean"   "area_mean"  
[5] "smoothness_mean"  "compactness_mean"  
[7] "concavity_mean"   "concave.points_mean"  
[9] "symmetry_mean"    "fractal_dimension_mean"  
[11] "radius_se"        "texture_se"  
[13] "perimeter_se"     "area_se"  
[15] "smoothness_se"    "compactness_se"  
[17] "concavity_se"     "concave.points_se"  
[19] "symmetry_se"      "fractal_dimension_se"  
[21] "radius_worst"     "texture_worst"  
[23] "perimeter_worst"  "area_worst"  
[25] "smoothness_worst" "compactness_worst"  
[27] "concavity_worst"  "concave.points_worst"  
[29] "symmetry_worst"   "fractal_dimension_worst"
```

```
length(grep("_mean", colnames(wisc.data), value = TRUE))
```

```
[1] 10
```

There are 10 variables in the data suffixed with `_mean`.

Principal Component Analysis (PCA)

We need to scale our data before PCA with the `scale=TRUE` argument to `prcomp()`.

```
wisc.pr <- prcomp(wisc.data, scale = TRUE)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal component (PC1)?

44.27% is captured by the first principal component.


Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

The first three principal components are required to describe at least 70% of the original variance in the data.

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

The first seven principal components are required to describe at least 90% of the original variance in the data.

```
biplot(wisc.pr)
```




Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

This plot is incredibly difficult to understand as conclusions cannot be made from the plot with everything crowded together.

```
library(ggplot2)


ggplot(wisc.pr$x) +
```

```
aes(PC1, PC2, col=diagnosis) +  
geom_point()
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
ggplot(wisc.pr$x) +  
aes(PC1, PC3, col=diagnosis) +  
geom_point()
```



For both plots, the malignant observations are towards the left of the graph whereas the benign are assembled towards the right. This presents a separation between the benign and malignant samples.


Variance explained

```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```


```
[1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```

```
# Variance explained by each principal component: pve
pve <- pr.var / sum(pr.var)

# Plot variance explained for each principal component
plot(c(1,pve), xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```



```
# Alternative scree plot of the same data, note data driven y-axis
barplot(pve, ylab = "Percent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```


```
## ggplot based graph
#install.packages("factoextra")
library(factoextra)
```

Warning: package 'factoextra' was built under R version 4.5.2

Welcome to factoextra!

Want to learn more? See two factoextra-related books at <https://www.datanovia.com/en/product/guide-to-principal-component-methods-in-r/>

```
fviz_eig(wisc.pr, addlabels = TRUE)
```



Communicating PCA results

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC. Are there any features with larger contributions than this one?

```
wisc.pr$rotation["concave.points_mean", 1]
```


```
[1] -0.2608538
```

The loading of `concave.points_mean` on PC1 is -0.2609 , which indicates a strong contribution to the first principal component. There are other features with larger contributions, such as `radius_mean`.

Hierarchical clustering

```
# Scale the wisc.data data using the "scale()" function
# Next, calculate the (Euclidean) distances between all pairs of observations in the new scaled data
wisc.hclust <- hclust(dist(scale(wisc.data)))
plot(wisc.hclust)
```

Cluster Dendrogram




```
dist(scale(wisc.data))
hclust (*, "complete")
```

Q10. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters?

```
plot(wisc.hclust)
abline(h = 19, col = "red", lty = 2)
```

Cluster Dendrogram



```
dist(scale(wisc.data))
hclust(*, "complete")
```

```
wisc.hclust.clusters <- cutree(wisc.hclust, k = 4)
table(wisc.hclust.clusters, diagnosis)
```

	diagnosis	
wisc.hclust.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2


Q12. Which method gives your favorite results for the same data.dist dataset?
Explain your reasoning.

The ward.D2 method gives the best results for this dataset. It produces more clearly separated clusters compared to methods like “single”, “complete”, and “average”. Because ward.D2 minimizes variance within clusters, it forms spherical clusters, which aligns well with the structure of this data and results in better separation of malignant and benign samples.

Combining methods

```
d <- dist(wisc.pr$x[,1:7])
wisc.pr.hclust <- hclust(d, method = "ward.D2")
plot(wisc.pr.hclust)
abline(h = 70, col = "red")
```

Cluster Dendrogram



d
hclust (*, "ward.D2")

```
grps <- cutree(wisc.pr.hclust, k = 2)
table(grps)
```


```
grps
  1  2
216 353
```

How does this clustering `grps` correspond to the expert diagnosis?

```
table(grps, diagnosis)
```

```
      diagnosis
grps   B   M
  1   28 188
  2  329  24
```

```
ggplot(wisc.pr$x) +
  aes(PC1, PC2) +
  geom_point(col=grps)
```



```
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
```

Q13. How well does the newly created hclust model with two clusters separate out the two “M” and “B” diagnoses?

```
table(wisc.pr.hclust.clusters, diagnosis)
```

	diagnosis	
wisc.pr.hclust.clusters	B	M
1	28	188
2	329	24

The newly created hclust model with two clusters separate out the two diagnoses pretty well with only 52 miscalculations. This shows an approximately 91% accuracy.

Q14. How well do the hierarchical clustering models you created in the previous sections (i.e. without first doing PCA) do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.hclust.clusters and wisc.pr.hclust.clusters) with the vector containing the actual diagnoses.

```
table(wisc.hclust.clusters, diagnosis)
```

```

              diagnosis
wisc.hclust.clusters  B   M
1    12 165
2     2   5
3   343  40
4     0   2

```

The hierarchical clustering model without PCA performs worse at separating malignant and benign samples. The clusters are more fragmented and show greater mixing of diagnoses.

Prediction

```

#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc

```

```


      PC1      PC2      PC3      PC4      PC5      PC6      PC7
[1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
[2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
      PC8      PC9      PC10     PC11     PC12     PC13     PC14
[1,] -0.2307350 0.1029569 -0.9272861 0.3411457  0.375921 0.1610764 1.187882
[2,] -0.3307423 0.5281896 -0.4855301 0.7173233 -1.185917 0.5893856 0.303029
      PC15     PC16     PC17     PC18     PC19     PC20
[1,] 0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
[2,] 0.1299153  0.1448061 -0.40509706  0.06565549  0.25591230 -0.4289500
      PC21     PC22     PC23     PC24     PC25     PC26
[1,] 0.1228233 0.09358453 0.08347651 0.1223396 0.02124121 0.078884581
[2,] -0.1224776 0.01732146 0.06316631 -0.2338618 -0.20755948 -0.009833238
      PC27     PC28     PC29     PC30
[1,] 0.220199544 -0.02946023 -0.015620933 0.005269029
[2,] -0.001134152 0.09638361 0.002795349 -0.019015820

```

```

plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")

```



Q16. Which of these new patients should we prioritize for follow up based on your results?

Patient 1 should be prioritized for follow up based on these results, as their PCA projection is closer to the malignant cluster, while Patient 2 appears more consistent with benign samples.

About this document

```
sessionInfo()
```

```
R version 4.5.1 (2025-06-13)
Platform: x86_64-apple-darwin20
Running under: macOS Sequoia 15.4.1
```

```
Matrix products: default
```

```
BLAS:   /Library/Frameworks/R.framework/Versions/4.5-x86_64/Resources/lib/libRblas.0.dylib
LAPACK: /Library/Frameworks/R.framework/Versions/4.5-x86_64/Resources/lib/libRlapack.dylib;
```

```
locale:
```

```
[1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8
```


time zone: America/Los_Angeles
tzcode source: internal

attached base packages:

[1] stats graphics grDevices utils datasets methods base

other attached packages:

[1] factoextra_2.0.0 ggplot2_4.0.2

loaded via a namespace (and not attached):

[1] gtable_0.3.6	jsonlite_2.0.0	dplyr_1.2.0	compiler_4.5.1
[5] ggsignif_0.6.4	tidyselect_1.2.1	Rcpp_1.1.0	tidyr_1.3.1
[9] scales_1.4.0	yaml_2.3.10	fastmap_1.2.0	R6_2.6.1
[13] ggpubr_0.6.3	labeling_0.4.3	generics_0.1.4	Formula_1.2-5
[17] knitr_1.50	backports_1.5.0	ggrepel_0.9.6	tibble_3.3.0
[21] car_3.1-3	pillar_1.11.1	RColorBrewer_1.1-3	rlang_1.1.7
[25] broom_1.0.10	xfun_0.53	S7_0.2.0	cli_3.6.5
[29] withr_3.0.2	magrittr_2.0.4	digest_0.6.37	grid_4.5.1
[33] rstudioapi_0.17.1	lifecycle_1.0.5	vctrs_0.7.1	rstatix_0.7.3
[37] evaluate_1.0.5	glue_1.8.0	farver_2.1.2	abind_1.4-8
[41] carData_3.0-5	rmarkdown_2.30	purrr_1.1.0	tools_4.5.1
[45] pkgconfig_2.0.3	htmltools_0.5.8.1		